Expanding-window walk-forward forecasting protocol

365 Fridays − 3 leading weeks − 1 trailing week = 361 eligible sequences = 104 warm-up + 257 OOS origins

Common outer calendar with expanding estimation folds and periodic re-estimation

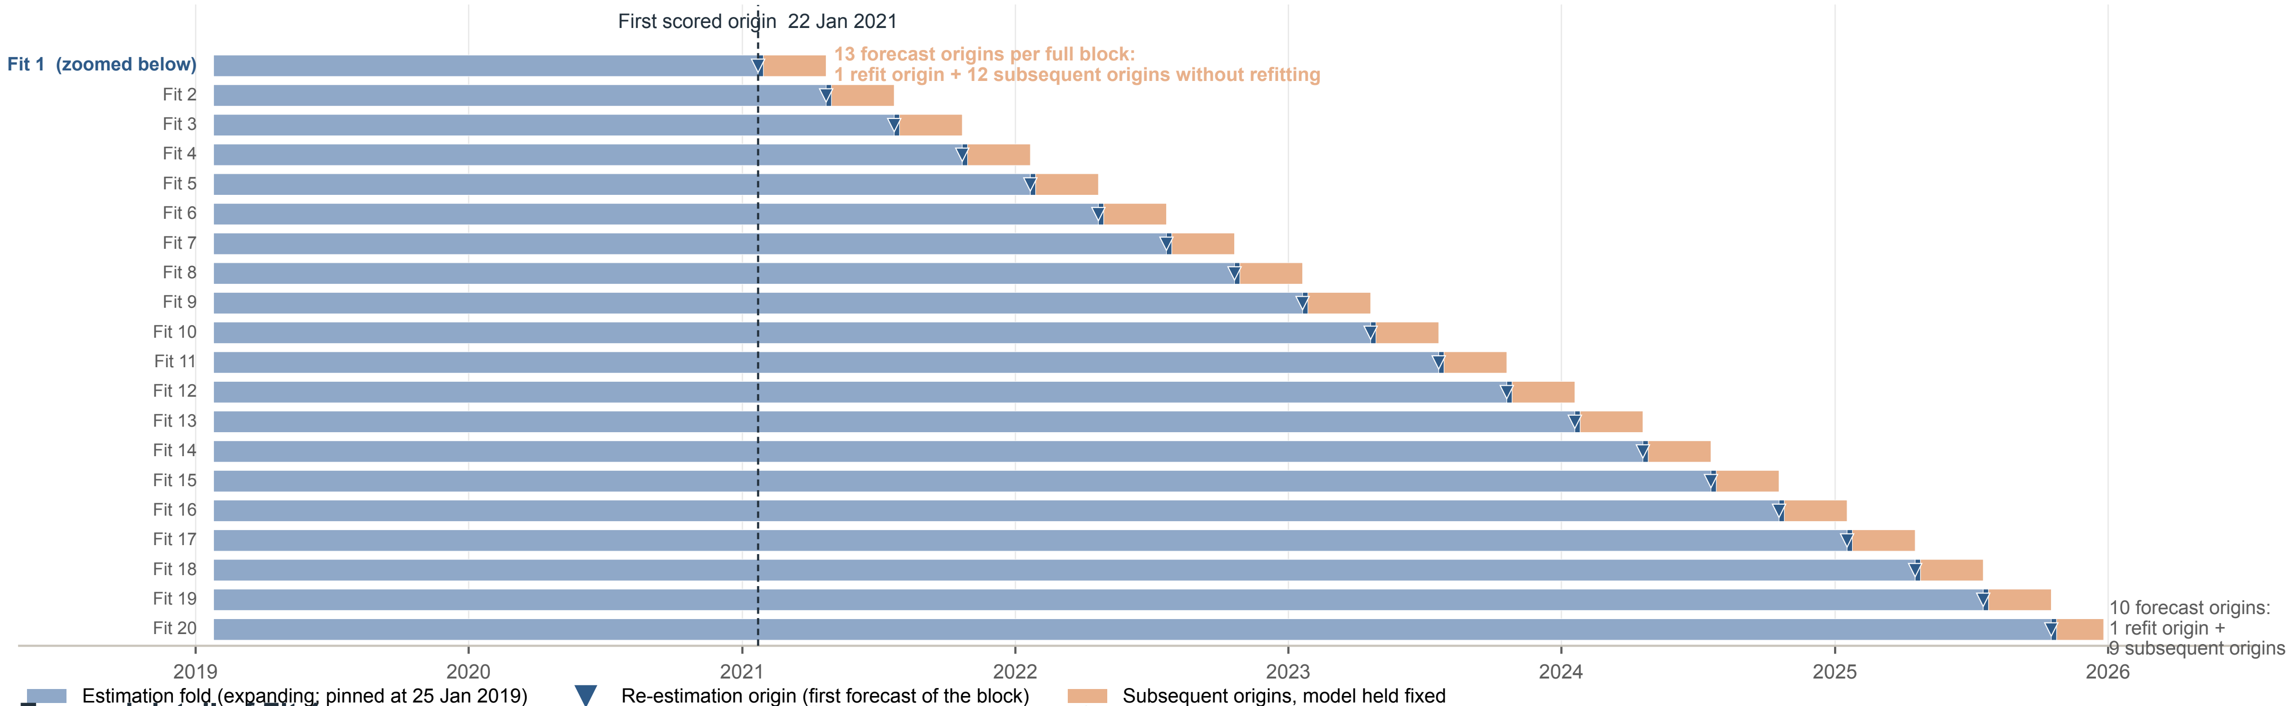

Zoomed detail of Fit 1

Different inner-validation rules for Flat and Deep

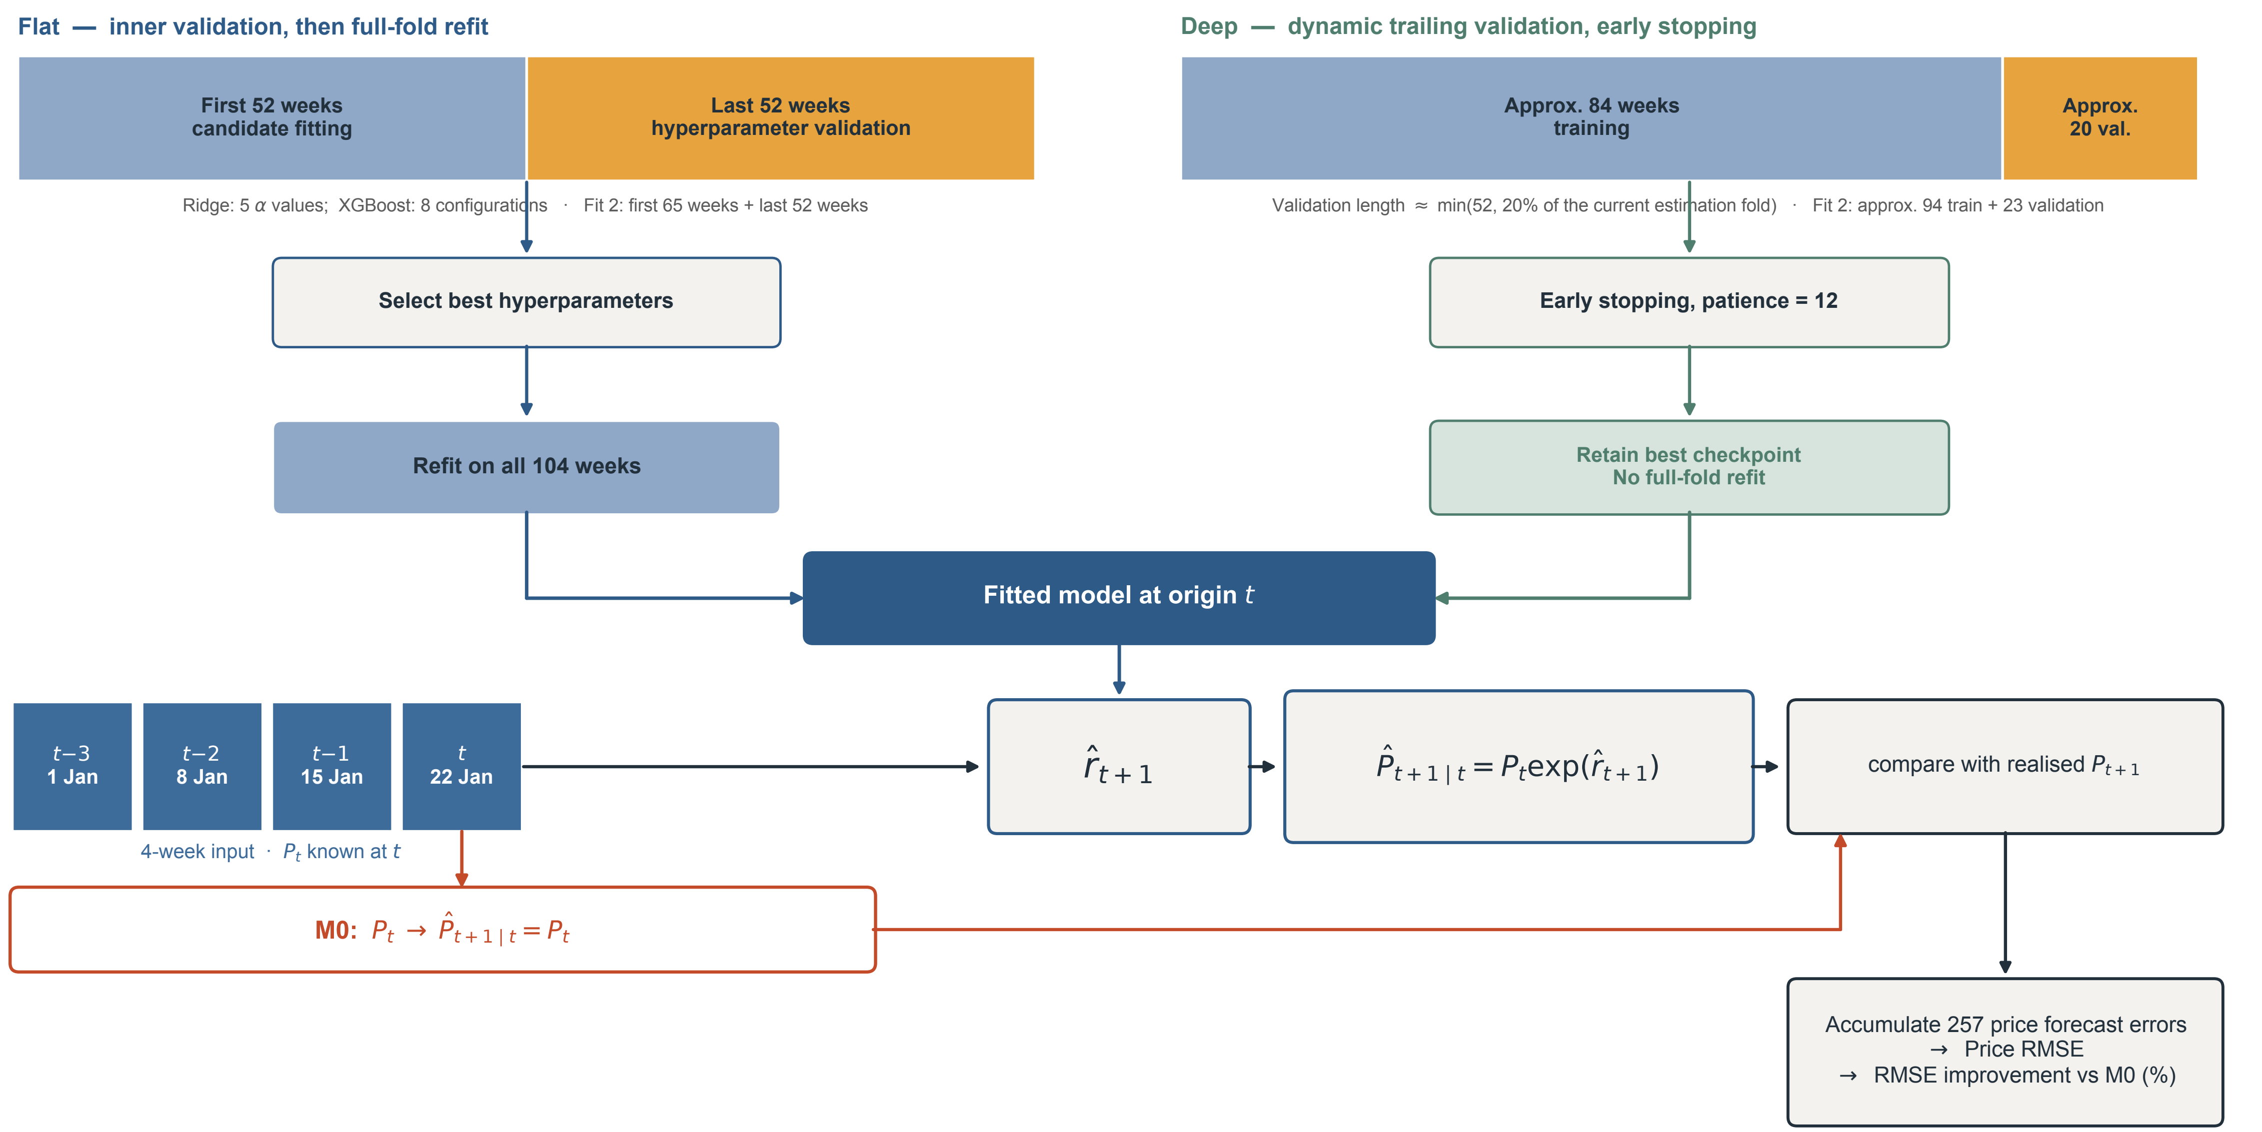

The estimation fold expands from an initial 104 sequences and is re-estimated every 13 forecast origins. Flat models use the final 52 training sequences for hyperparameter validation and are subsequently refit on the full estimation fold. Deep models use a dynamically sized trailing validation segment for early stopping and retain the best checkpoint without full-fold retraining. Both model families are evaluated on the same 257 out-of-sample forecast origins.